


Gradient Boosting Machines



Amy Ruskin
MUMT 621



A note on terminology

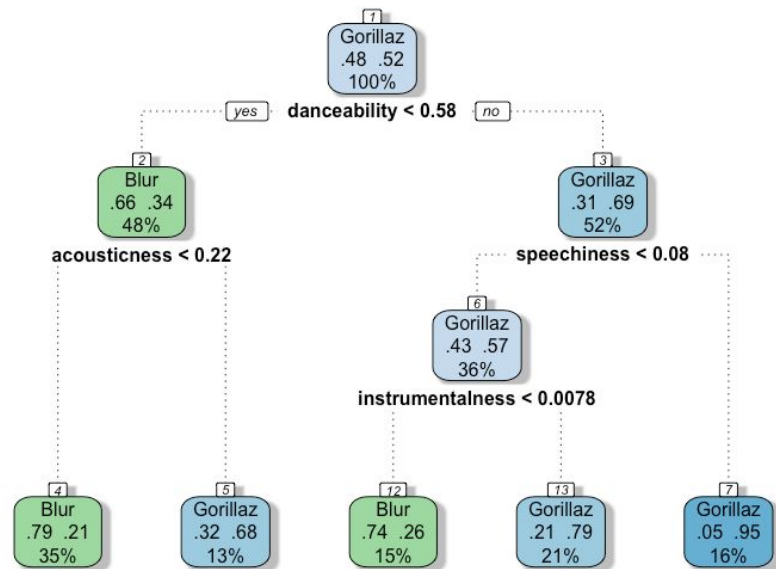
- **Gradient boosting machine (GBM)** (Friedman 2001)
- Gradient boosting decision trees (GBDT) (Ke et al. 2017; Dorogush et al. 2017; Bai and Fan 2017)
- Gradient tree boosting (Chen and Guestrin 2016)
- Gradient boosted regression tree (Chen and Guestrin 2016)
- Gradient-boosted trees (GBTs)¹

¹<https://spark.apache.org/docs/latest/mllib-ensembles.html#gradient-boosted-trees-gbts>

Decision trees

- Partition feature space into rectangles
- Observations in same partition assigned same value
- At each level, find variable and split point to minimize splitting criterion across left and right children
- Examples of splitting criteria
 - **Regression:** sum of squared residuals
 - **Classification:** measures of node impurity like misclassification rate, Gini index, cross-entropy/deviance

(Hastie et al. 2009)



A simple example of a decision tree for artist classification (using Spotify audio features)

Pros & cons of decision trees

Advantages

- Fast to construct
- Easy to interpret models
- Mix numeric and categorical predictors
- Handling of missing values
- Invariant under strictly monotone transformations of predictors
 - Not sensitive to long-tailed distributions and outliers in predictors
- Feature selection performed as part of fitting procedure

Disadvantages

- Instability—highly susceptible to small changes in training data
- Lack of smoothness in prediction surface
- Difficulty capturing additive structure
- Inaccurate compared to other methods

(Hastie et al. 2009;
Friedman 2001)

TABLE 10.1. *Some characteristics of different learning methods. Key: ▲ = good, ◆ = fair, and ▼ = poor.*

Characteristic	Neural Nets	SVM	Trees	MARS	k-NN, Kernels
Natural handling of data of “mixed” type	▼	▼	▲	▲	▼
Handling of missing values	▼	▼	▲	▲	▲
Robustness to outliers in input space	▼	▼	▲	▼	▲
Insensitive to monotone transformations of inputs	▼	▼	▲	▼	▼
Computational scalability (large N)	▼	▼	▲	▲	▼
Ability to deal with irrelevant inputs	▼	▼	▲	▲	▼
Ability to extract linear combinations of features	▲	▲	▼	▼	◆
Interpretability	▼	▼	◆	▲	▼
Predictive power	▲	▲	▼	◆	▲

Boosting can improve **accuracy** of trees at the cost of **interpretability** and **speed**

(Hastie et al. 2009, p. 351)

What is boosting?

- Combination of “weak” models → powerful “committee” (Hastie et al. 2009)
- Iteratively fit models to compensate for shortcomings in previously trained models (Li 2016)

Gradient tree boosting algorithm

$L(y, f(x))$: some differentiable loss function

M : total # of trees in model

N : total # of observations in training data

(Hastie et al. 2009, p. 361)

1. Initialize $f_0(x) = \arg \min_{\gamma} \sum_{i=1}^N L(y_i, \gamma)$.

2. For $m = 1$ to M :

(a) For $i = 1, 2, \dots, N$ compute

$$r_{im} = - \left[\frac{\partial L(y_i, f(x_i))}{\partial f(x_i)} \right]_{f=f_{m-1}}.$$

Pseudo-residual (for observation i , based on predictions from trees 1 to $m-1$)

(b) Fit a regression tree to the targets r_{im} giving terminal regions R_{jm} , $j = 1, 2, \dots, J_m$.

(c) For $j = 1, 2, \dots, J_m$ compute

$$\gamma_{jm} = \arg \min_{\gamma} \sum_{x_i \in R_{jm}} L(y_i, f_{m-1}(x_i) + \gamma).$$

Value assigned at leaf j in tree m



(d) Update $f_m(x) = f_{m-1}(x) + \sum_{j=1}^{J_m} \gamma_{jm} I(x \in R_{jm})$.

3. Output $\hat{f}(x) = f_M(x)$.


GBM parameters

- Number of trees in model
- Shrinkage
 - Scale each update by “learning rate” parameter
 - Lower learning rate → more trees needed to converge
- Size of constituent trees
 - Constrain maximum depth, minimum number of observations per leaf
- Penalization
- Bagging/subsampling

Example: Making predictions with a GBM



...



		Predicted P(Artist = Blur)			
Artist	Song	1 Tree	2 Trees	...	50 Trees
Blur	"Tracy Jacks"	0.509	0.489		0.619
Blur	"Beetlebum"	0.509	0.526		0.759
Gorillaz	"DARE"	0.468	0.448		0.223
Gorillaz	"Feel Good Inc."	0.468	0.448		0.135

Tree 0:

Danceability: $0.636 < 0.663$
 $P(\text{Blur}) = \text{logit}^{-1}(0.037) = 0.509$

Tree 1:

Danceability: $0.636 > 0.550$
 $P(\text{Blur}) = \text{logit}^{-1}(0.037 - 0.079) = 0.489$

...

Tree 49:

Acousticness: $0.098 > 0.080$
 $P(\text{Blur}) = \text{logit}^{-1}(0.037 - 0.079 + \dots - 0.038) = 0.619$

Pros & cons of GBMs

Advantages

- Over single decision tree:
 - Improved accuracy
 - Prediction surface still not smooth, but finer granularity to piecewise approximations
 - Averaging over many small trees → greater stability
- Maintains many desirable properties of decision trees
- Can give quick indication of dataset's potential predictability

Disadvantages

- Slower to compute than single decision tree
- Difficult to interpret model
 - Relative importance of input variables
 - Partial dependence plots

(Hastie et al. 2009;
Friedman 2001)

Implementations

- Hot right now:
 - XGBoost (<https://github.com/dmlc/xgboost>)
 - Developed at University of Washington, first release in 2014
 - Official support in: Python, R, JVM, Ruby, Julia, C
 - LightGBM (<https://github.com/microsoft/LightGBM>)
 - Developed by Microsoft Research, first release in 2017
 - Official support in: Python, R, C
 - CatBoost (<https://github.com/catboost>)
 - Developed at Yandex, first release in 2017
 - Official support in: Python, R (apply already trained models in Java, C/C++)
 - Learning tasks: regression, classification (binary and multiclass), ranking
- In machine learning libraries: h2o, scikit-learn, Spark MLlib, and more

Boosted tree methods for MIR

- “Aggregate Features and AdaBoost for Music Classification” (Bergstra et al. 2006)
 - Winning method for genre classification in MIREX 2005 contest, runner-up for artist recognition
 - AdaBoost classifier with small decision trees (or stumps) used as weak classifiers
- “Incorporating Field-aware Deep Embedding Networks and Gradient Boosting Decision Trees for Music Recommendation” (Bai and Fan 2017)
 - Winning entry to WSDM-KKBOX Music Recommendation Challenge using ensemble of Field-aware Deep Embedding Networks and GBMs (LightGBM)
- “Evaluating Music Mastering Quality Using Machine Learning” (Shtern et al. 2018)
 - Used CatBoost implementation of gradient boosting
- “Detecting Music Genre Using Extreme Gradient Boosting” (Murauer and Specht 2018)
 - Submission to crowdAI music genre classification challenge in WebConference 2018
 - Best results from XGBoost classifier trained on numeric features extracted from mp3 files with Essentia

References

- Bai, Bing, and Yushun Fan. 2017. “Incorporating Field-aware Deep Embedding Networks and Gradient Boosting Decision Trees for Music Recommendation.” In *The 11th ACM International Conference on Web Search and Data Mining (WSDM)*.
- Bergstra, James, Norman Casagrande, Dumitru Erhan, Douglas Eck, and Balázs Kégl. 2006. “Aggregate Features and AdaBoost for Music Classification.” *Machine Learning*, 65: 473–484. <https://doi.org/10.1007/s10994-006-9019-7>
- Chen, Tianqi, and Carlos Guestrin. 2016. “XGBoost: A Scalable Tree Boosting System.” In *Proceedings of the 22nd ACM SIGKDD International Conference on Knowledge Discovery and Data Mining*, 785–794. <https://doi.org/10.1145/2939672.2939785>
- Dorogush, Anna Veronika, Vasily Ershov, and Andrey Gulin. 2018. “CatBoost: Gradient Boosting with Categorical Features Support.” *arXiv preprint arXiv:1810.11363*
- Friedman, Jerome H. 2001. “Greedy Function Approximation: A Gradient Boosting Machine.” *The Annals of Statistics* 29 (5): 1189–1232.

References

- Hastie, Trevor, Robert Tibshirani, and Jerome Friedman. 2009. *The Elements of Statistical Learning* (2nd ed). New York: Springer.
- Ke, Guolin, Qi Meng, Thomas Finley, Taifeng Wang, Wei Chen, Weidong Ma, Qiwei Ye, Tie-Yan Liu. 2017. “LightGBM: A Highly Efficient Gradient Boosting Decision Tree.” *Advances in Neural Information Processing Systems* 30 (NIPS 2017): 3149–3157.
- Li, Cheng. 2016. “A Gentle Introduction to Gradient Boosting.” https://www.chengli.io/tutorials/gradient_boosting.pdf (accessed 17 February 2020)
- Murauer, Benjamin, and Günther Specht. 2018. “Detecting Music Genre Using Extreme Gradient Boosting.” In *Companion Proceedings of the The Web Conference 2018*, 1923–27. <https://doi.org/10.1145/3184558.3191822>
- Shtern, Mark, Pedro Casas, and Vassilios Tzerpos. 2018. “Evaluating Music Mastering Quality Using Machine Learning.” In *Proceedings of the 28th Annual International Conference on Computer Science and Software Engineering*, 126–35.